

# Workpackage 5: High Performance Mathematical Computing

Clément Pernet

Final OpenDreamKit Project review

Luxembourg, October 30, 2019

## Mathematical computing

Computing with a large variety of objects

►  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

►  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

► Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

- ▶  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$
$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

- ▶  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices of polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

$$\begin{bmatrix} 3x^2+3 & 2x^2+3 \\ 4x^2+1 & x^2+4x+4 \end{bmatrix}$$

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

- ▶  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices of polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

$$\begin{bmatrix} 3x^2+3 & 2x^2+3 \\ 4x^2+1 & x^2+4x+4 \end{bmatrix}$$

- ▶ Tree algebras

for applications where all digits matter.

$$\frac{3q^2 - q^5}{q^5 + 2q^4 + 3q^3 + 3q^2 + 2q + 1} \begin{array}{c} (a) \\ / \quad \backslash \\ (b) \quad (c) \\ \quad | \\ \quad (d) \end{array} + \frac{2q}{q^4 + q^3 + 2q^2 + q + 1} \begin{array}{c} (a) \\ | \\ (b) \\ / \quad \backslash \\ (c) \quad (d) \end{array}$$

# High performance mathematical computing

**Need for High performance:** applications where size is crucial:

**Experimental maths:** testing conjectures

- larger instances give higher confidence

# High performance mathematical computing

**Need for High performance:** applications where size is crucial:

**Experimental maths:** testing conjectures

- ▶ larger instances give higher confidence

**Algebraic cryptanalysis:** security = computational difficulty

- ▶ key size determined by the largest solvable problem

## Example

Breaking RSA by integer factorization:  $n = pq$ . Last record:

- ▶  $n$  of 768 bits
- ▶ linear algebra in dimension 192 796 550 over  $\mathbb{F}_2$  (105Gb)
- ▶ About 2000 CPU years



# High performance mathematical computing

**Need for High performance:** applications where size is crucial:

**Experimental maths:** testing conjectures

- ▶ larger instances give higher confidence

**Algebraic cryptanalysis:** security = computational difficulty

- ▶ key size determined by the largest solvable problem

## Example

Breaking RSA by integer factorization:  $n = pq$ . Last record:

- ▶  $n$  of 768 bits
- ▶ linear algebra in dimension 192 796 550 over  $\mathbb{F}_2$  (105Gb)
- ▶ About 2000 CPU years

**3D data analysis, shape recognition:**

- ▶ via persistent homology
- ▶ large sparse matrices over  $\mathbb{F}_2, \mathbb{Z}$

# Goal: delivering high performance to maths users

**Systems :**

**GAP**

**PARI/GP**

**SageMath**

**Singular**

**Components :**

**FLINT**

**MPIR**

**LinBox**

**PPL**

**NumPy**

# Goal: delivering high performance to maths users

## Harnessing modern hardware $\rightsquigarrow$ parallelisation

- ▶ in-core parallelism (SIMD vectorisation)
- ▶ multi-core parallelism
- ▶ distributed computing: clusters, cloud

**Systems :**

**GAP**

**PARI/GP**

**SageMath**

**Singular**

**Components :**

**FLINT**

**MPIR**

**LinBox**

**PPL**

**NumPy**

**Architectures :**

**SIMD**

**Multicore  
server**

**HPC cluster**

**Cloud**

# Goal: delivering high performance to maths users

## Languages

- ▶ Computational Maths software uses high level languages (e.g. Python)
  - ▶ High performance delivered by languages close to the metal (C, assembly)
- ~> compilation, automated optimisation

**Systems :**

**GAP**

**PARI/GP**

**SageMath**

**Singular**

**Components :**

**FLINT**

**MPIR**

**LinBox**

**PPL**

**NumPy**

**Languages :**

**Cython**

**Python**

**Pythran**

**C**

**Architectures :**

**SIMD**

**Multicore  
server**

**HPC cluster**

**Cloud**

## Goal:

- ▶ Improve/Develop parallelization of software components
- ▶ Expose them through the software stack
- ▶ Offer High Performance Computing to VRE's users

# Outcome of WorkPackage 5

| Component           | Review 1   | Review 2 | Final review |
|---------------------|------------|----------|--------------|
| T5.1 Pari/GP        |            |          | D5.16        |
| T5.2 GAP            |            |          | D5.15        |
| T5.3 LinBox         |            | D5.12    | D5.14        |
| T5.4 Singular       | D5.6, D5.7 |          | D5.13        |
| T5.5 MPIR           | D5.5, D5.7 |          |              |
| T5.6 Combinatorics  | D5.1       | D5.11    |              |
| T5.7 Pythran        | D5.2       | D5.11    |              |
| T5.8 SunGrid Engine | D5.3       |          |              |

## Overall

- ▶ 31 software releases
- ▶ 16 research papers

# Outline

## Context and Focus

T5.1: Number theory with PARI/GP

T5.2: Group theory with GAP

T5.3: Exact linear algebra with LinBox

T5.4: Polynomial arithmetic with Singular



# Outline


## Context and Focus

T5.1: Number theory with PARI/GP

T5.2: Group theory with GAP

T5.3: Exact linear algebra with LinBox

T5.4: Polynomial arithmetic with Singular





# Computational Kernels: linear algebra

*Mathematics is the art of reducing any problem to linear algebra*

— W. Stein

# Computational Kernels: linear algebra

*Mathematics is the art of reducing any problem to linear algebra*  
— W. Stein

**Linear algebra: a key building block for HPC**

## Similarities with numerical HPC

- ▶ central elementary problem to which others reduce to
- ▶ (rather) simple algorithmic
- ▶ high compute/memory intensity

# Computational Kernels: linear algebra

*Mathematics is the art of reducing any problem to linear algebra*

SINGULAR – W. Stein

**Linear algebra: a key building block for HPC**

## Similarities with numerical HPC

- ▶ central elementary problem to which others reduce to
- ▶ (rather) simple algorithmic
- ▶ high compute/memory intensity

## Specificities

- ▶ Multiprecision arithmetic  $\Rightarrow$  lifting from finite precision ( $\mathbb{F}_p$ )
- ▶ Rank deficiency  $\Rightarrow$  unbalanced dynamic blocking
- ▶ Early adopter of subcubic matrix arithmetic  $\Rightarrow$  recursion

# Computational kernels: arithmetic

- ▶ Wide variety of computing domains:  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q, \mathbb{F}_q[X], \mathbb{Z}[X, Y, Z], \dots$
- ▶ possibly with dynamic size

## Challenge

- ▶ most often memory intensive operations → hard to parallelize
- ▶ very fine grain, but billions of instance → fine tuning

# Outline

## Context and Focus

T5.1: Number theory with PARI/GP

T5.2: Group theory with GAP

T5.3: Exact linear algebra with LinBox

T5.4: Polynomial arithmetic with Singular



## PARI ecosystem

**PARI library:** dedicated routines for number theory

**PARI-GP:** an interactive system

**GP2C:** a GP to C compiler

## Generic parallelization engine for the whole suite

Delivers support for

Features:

- ▶ sequential computations
- ▶ POSIX threads
- ▶ MPI for distributed computing
- ▶ Same code base
- ▶ automated parallelization
- ▶ full control for power users/developpers

# Main Achievements

- ▶ Fast linear algebra over the rationals and cyclotomic fields
- ▶ Fast Chinese remainders and multimodular reduction
- ▶ Parallel polynomial resultant
- ▶ Fast modular polynomials and applications
- ▶ MPQS integer factorization rewrite

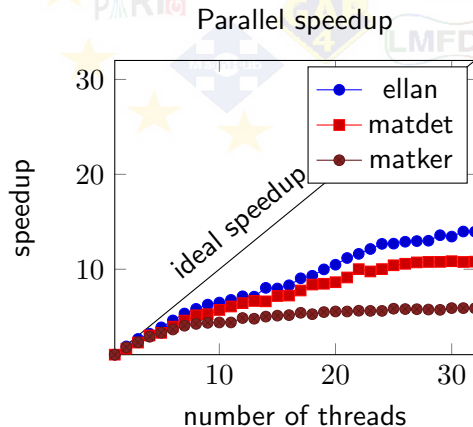
## Well-honed strategy after preliminary assessment

- ▶ Creation of "worker" functions from existing code
- ▶ Insertion of actual parallel instructions
- ▶ Incremental buildup, independently instrumenting one high-level function at a time.

Linear algebra over rationals:

- ▶ Required for cyclotomic rings (Allombert, 2016-17)
- ▶ Fast Chinese remaindering (Allombert, 2017)
- ▶ Fast CUP decomposition over finite fields (Bruin, 2018)
- ▶ Parallelization (Allombert, Belabas, 2017-19)

Fourier transform of  $L$ -functions





# Outline

## Context and Focus

T5.1: Number theory with PARI/GP

T5.2: Group theory with GAP

T5.3: Exact linear algebra with LinBox

T5.4: Polynomial arithmetic with Singular




## HPC-GAP

- ▶ Multi-threaded GAP.
- ▶ Targets:
  - ▶ multicore servers.
  - ▶ good speedups
  - ▶ high level abstraction

## Achievement

- ▶ Fork from GAP, diverged for a long time
- ▶ Huge effort to bring it back in:  
GAP 4.9.1 (Month 33).
  - first GAP release with HPCGAP integrated as a compile-time option



# Dense linear algebra over small finite fields

- ▶ Matrix multiplication, Gaussian elimination, echelon forms, etc
- ▶ A key kernel for many GAP computations

## MeatAxe64

A new C and assembler library, tuned for performance at all levels:

- ▶ new data representations and assembler kernels
- ▶ new algorithms for many fields
- ▶ control of cache usage and memory bandwidth
  - allowing for sharing between threads and cores
- ▶ purpose built highly efficient task farm
  - $1\text{M} \times 1\text{M}$  dense matrix multiply over  $\text{GF}(2)$  in 8 hours (64 core AMD bulldozer).

Fully available from GAP.

# Outline

## Context and Focus

T5.1: Number theory with PARI/GP

T5.2: Group theory with GAP

T5.3: Exact linear algebra with LinBox

T5.4: Polynomial arithmetic with Singular



# Parallel Rational solver algorithmic

| Method                  | Bit complexity   |
|-------------------------|------------------|
| Gauss over $\mathbb{Q}$ | $2^{O(n)}$       |
| Gauss mod bound(sol)    | $O(n^5)$         |
| Chinese Remaindering    | $\tilde{O}(n^4)$ |
| $p$ -adic lifting       | $\tilde{O}(n^3)$ |



## Chinese Remainder

1. Solve the system independently modulo  $p_1, p_2, \dots, p_k$
2. Reconstruct a solution modulo  $p_1 \times p_2 \times \dots, p_k$ .
3. Reconstruct over  $\mathbb{Q}$

## $p$ -adic lifting

1. Solve the system modulo  $p$
2. Iteratively lift the solution modulo  $p^2, p^3, \dots, p^k$
3. Reconstruct over  $\mathbb{Q}$

# Distributed memory Chinese Remaindering



## Conclusions

- ▶ (almost) embarrassingly parallel
- ▶ but overwhelming computational cost ( $O(n^4)$ )
- ▶ hybrid OpenMP-MPI version slightly slower but better memory efficiency

# Shared memory $p$ -adic lifting

## A new hybrid algorithm: Chinese Remaindering within $p$ -adic lifting

- ▶ Smaller critical path
- ▶ Higher degree of parallelism


# Shared memory $p$ -adic lifting

## A new hybrid algorithm: Chinese Remaindering within $p$ -adic lifting

- ▶ Smaller critical path
- ▶ Higher degree of parallelism


## Improving state of the art efficiency

- ▶ Improved sequential efficiency (improved memory access pattern, BLAS3)
- ▶ Chinese remaindering delivers good parallel scaling






# GPU enabled fflas-ffpack



- ▶ Matrix product over  $\mathbb{Z}/131071\mathbb{Z}$
- ▶ 32 Skylake-X cores (AVX512)
- ▶ 1 Nvidia Tesla V100 GPU

# GPU enabled fflas-ffpack



- ▶ Matrix product over  $\mathbb{Z}/131071\mathbb{Z}$
- ▶ 32 Skylake-X cores (AVX512)
- ▶ 1 Nvidia Tesla V100 GPU

## Limitations and perspectives

### Bottleneck in the transfer between GPU and RAM

- ▶ deport more computations to the GPU → dedicated GPU kernels
- ▶ communication avoiding block scheduling → deep structural change

# Outline

## Context and Focus

T5.1: Number theory with PARI/GP


T5.2: Group theory with GAP

T5.3: Exact linear algebra with LinBox

T5.4: Polynomial arithmetic with Singular




## Multivariate Polynomial Multiplication in FLINT




- ▶ over  $\mathbb{Z}$ ,  $\mathbb{Q}$ ,  $\mathbb{F}_p$ ,  $\mathbb{F}_{p^n}$  for  $p < 2^{64}$ .
- ▶ Lex, degLex and degGrevLex ordering supported
- ▶ Sequential alg: improves Singular's
- ▶ Close to linear scaling
- ▶ Singular now relies on FLINT

## Multivariate Polynomial Multiplication in FLINT




- ▶ over  $\mathbb{Z}$ ,  $\mathbb{Q}$ ,  $\mathbb{F}_p$ ,  $\mathbb{F}_{p^n}$  for  $p < 2^{64}$ .
- ▶ Lex, degLex and degGrevLex ordering supported
- ▶ Sequential alg: improves Singular's
- ▶ Close to linear scaling
- ▶ Singular now relies on FLINT

## Multivariate Polynomial GCD in FLINT



- ▶ over  $\mathbb{Z}$ ,  $\mathbb{Q}$ ,  $\mathbb{F}_p$ ,  $\mathbb{F}_{p^n}$  for  $p < 2^{64}$ .
- ▶ Lex, degLex and degGrevLex ordering supported
- ▶ Sequential alg: improves Singular's
- ▶ Close to linear scaling
- ▶ Singular now relies on FLINT

## Multivariate Polynomial GCD in FLINT



- ▶ over  $\mathbb{Z}$ ,  $\mathbb{Q}$ ,  $\mathbb{F}_p$ ,  $\mathbb{F}_{p^n}$  for  $p < 2^{64}$ .
- ▶ Lex, degLex and degGrevLex ordering supported
- ▶ Sequential alg: improves Singular's
- ▶ Close to linear scaling
- ▶ Singular now relies on FLINT

## Parallelizing memory intensive kernels

- ▶ Overhead of split and combine
- ▶ Only amortized for large instances
- ▶ Parallelization at a coarser grain



## Parallelizing memory intensive kernels

- ▶ Overhead of split and combine
- ▶ Only amortized for large instances
- ▶ Parallelization at a coarser grain

## Data placement

- ▶ Heavy use of dynamic allocation (`malloc`)
- ▶ Causes performances fickle
- ▶ To be used with parsimony
- ▶ Investigate dedicated allocators

# New Perspectives and directions

## Parallelizing memory intensive kernels

- ▶ Overhead of split and combine
- ▶ Only amortized for large instances
- ▶ Parallelization at a coarser grain

## Data placement

- ▶ Heavy use of dynamic allocation (`malloc`)
- ▶ Causes performances fickle
- ▶ To be used with parsimony
- ▶ Investigate dedicated allocators

## New directions

Application driven:

- ▶ More orderings: block, weighted
- ▶ Factorization: (harnessing most T5.4 contributions)

# Conclusion

## Outcome of WorkPackage 5

|                     | Review 1   | Review 2 | Final review |
|---------------------|------------|----------|--------------|
| T5.1 Pari/GP        |            |          | D5.16        |
| T5.2 GAP            |            |          | D5.15        |
| T5.3 LinBox         |            | D5.12    | D5.14        |
| T5.4 Singular       | D5.6, D5.7 |          | D5.13        |
| T5.5 MPIR           | D5.5, D5.7 |          |              |
| T5.6 Combinatorics  | D5.1       | D5.11    |              |
| T5.7 Pythran        | D5.2       | D5.11    |              |
| T5.8 SunGrid Engine | D5.3       |          |              |

### Overall

- ▶ 31 software releases
- ▶ 16 research papers

## From dedicated to general purpose HPC components:

- ▶ Early instances of HPC computer algebra: dedicated to some target application (breaking RSA, etc)
- ▶ Building a general purpose HPC component:
  - ▶ challenging
  - ▶ longer term sustainability

## Identifying the right place to focus efforts on

- ▶ Premature focus on embarrassingly parallel distributed computing
- ▶ Might as well improve a more difficult algorithm on multicore servers

## Risk of technology dependency

- ▶ Cilk: from success to shut-down
- ▶ Interchangeability and modularity (PARI, LinBox)

## Deployment of architecture dependent software stacks

- ▶ dynamic runtime configuration
  - user decision based on algorithms (high level), not systems (low level)
  - compilation hurdle, vs slick interface of the VRE
- ▶ extend high level control over parallel resources: GPUs, cluster nodes, etc
- ▶ keep as much HPC features as possible in virtualization images